


AI/ML Task

Develop a machine learning model to detect fraudulent transactions using a Kaggle dataset  
**Kaggle Credit Card Fraud Detection dataset** with a focus on data handling, model training, evaluation, and explainability.

Task 1: Data Exploration and Preprocessing

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
```


```
data = pd.read_csv("creditcard.csv")
data
```



	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0

284807 rows × 31 columns

```
print("Dataset Info: ", data.info())
print("\nClass Distribution: ", data['Class'].value_counts())
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Time      284807 non-null  float64
1    V1         284807 non-null  float64
2    V2         284807 non-null  float64
3    V3         284807 non-null  float64
4    V4         284807 non-null  float64
5    V5         284807 non-null  float64
6    V6         284807 non-null  float64
7    V7         284807 non-null  float64
8    V8         284807 non-null  float64
9    V9         284807 non-null  float64
10   V10        284807 non-null  float64
11   V11        284807 non-null  float64
12   V12        284807 non-null  float64
13   V13        284807 non-null  float64
14   V14        284807 non-null  float64
15   V15        284807 non-null  float64
16   V16        284807 non-null  float64
17   V17        284807 non-null  float64
18   V18        284807 non-null  float64
19   V19        284807 non-null  float64
20   V20        284807 non-null  float64
21   V21        284807 non-null  float64
22   V22        284807 non-null  float64
23   V23        284807 non-null  float64
```

```

24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount    284807 non-null float64
30 Class     284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
Dataset Info:  None

```

```

Class Distribution:  Class
0      284315
1        492
Name: count, dtype: int64

```

```

# Feature Scaling for 'Amount' (normalize)
scaler = StandardScaler()
data['Amount'] = scaler.fit_transform(data['Amount'].values.reshape(-1, 1))

# Split Features and Target
X = data.drop(columns=['Class'])
y = data['Class']

# Handle Class Imbalance using SMOTE
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

```

## ✓ Task 2: Supervised Model Development

```

# Importing necessary libraries for modeling
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve
from xgboost import XGBClassifier

# Logistic Regression

# Splitting Data into Train-Test Sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size=0.3, random_state=42)

# Training Logistic Regression (Baseline Model)
lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)
y_pred_lr = lr_model.predict(X_test)

# Evaluating Logistic Regression
print("Logistic Regression Performance:")
print(classification_report(y_test, y_pred_lr))

# Confusion Matrix for Logistic Regression
conf_matrix = confusion_matrix(y_test, y_pred_lr)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title("Logistic Regression - Confusion Matrix")
plt.show()

```

→ /usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>


Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Logistic Regression Performance:

	precision	recall	f1-score	support
0	0.97	0.98	0.97	85149
1	0.98	0.97	0.97	85440
accuracy			0.97	170589
macro avg	0.97	0.97	0.97	170589
weighted avg	0.97	0.97	0.97	170589



# XGBoost Model

# Training XGBoost Model

```
xgb_model = XGBClassifier(eval_metric='logloss', use_label_encoder=False)
xgb_model.fit(X_train, y_train)
y_pred_xgb = xgb_model.predict(X_test)
```

# Evaluating XGBoost Model

```
print("XGBoost Performance:")
print(classification_report(y_test, y_pred_xgb))
```

# Confusion Matrix for XGBoost

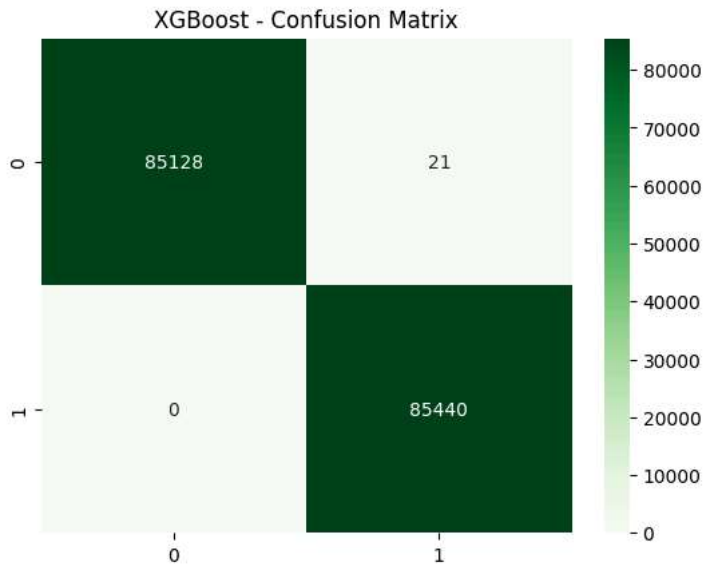
```
conf_matrix = confusion_matrix(y_test, y_pred_xgb)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Greens')
plt.title("XGBoost - Confusion Matrix")
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning: [14:52:03] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(msg, UserWarning)
XGBoost Performance:
      precision    recall  f1-score   support


      0       1.00      1.00      1.00     85149
      1       1.00      1.00      1.00     85440

 accuracy         1.00
 macro avg       1.00      1.00      1.00     170589
 weighted avg    1.00      1.00      1.00     170589
```



```
# ROC-AUC Curve for XGBoost
y_pred_prob = xgb_model.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
roc_auc = roc_auc_score(y_test, y_pred_prob)

plt.plot(fpr, tpr, label=f"XGBoost (AUC = {roc_auc:.2f})")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC-AUC Curve")
plt.legend()
plt.show()
```



### ✓ Task 3: Unsupervised Model Development

```
from sklearn.ensemble import IsolationForest
```

```
# Training Isolation Forest Model
isolation_forest = IsolationForest(contamination=0.01, random_state=42)
data['Anomaly'] = isolation_forest.fit_predict(X)
data['Anomaly'] = data['Anomaly'].apply(lambda x: 1 if x == -1 else 0)
```

```
# Evaluating Anomalies
anomalies_detected = data[data['Anomaly'] == 1]
print(f"Total Anomalies Detected: {len(anomalies_detected)}")
print("Sample Anomalies:", anomalies_detected.head())
```

```
→ Total Anomalies Detected: 2849
Sample Anomalies:
   Time  V1  V2  V3  V4  V5  V6  \
164 103.0 -6.093248 -12.114213 -5.694973 3.294389 -1.413792 4.776000
362 266.0 -2.564961 2.470985 2.649417 -1.564256 1.794297 -0.614742
401 290.0 -5.166299 -5.449369 2.988498 2.658991 1.948152 -0.854470
455 333.0 -2.977214 0.781748 2.881724 -1.627798 -1.368067 1.656876
601 454.0 -3.133891 2.887394 2.178690 -1.576535 1.434510 -0.687313

   V7  V8  V9  ...  V22  V23  V24  \
164 4.808426 -0.228197 -0.525896 ... -2.264037 -3.553381 1.215279
362 4.185906 -3.855359 5.436633 ... -0.463149 -0.532466 0.306494
401 -0.326394 -1.017364 1.983901 ... 1.172278 3.150413 0.574081
455 -1.185053 -5.393736 1.719407 ... -1.929181 0.190843 0.533134
601 3.816056 -3.416915 5.459274 ... -0.498518 -0.476668 0.302196

   V25  V26  V27  V28  Amount  Class  Anomaly
164 -0.406073 -0.653384 -0.711222 0.672900 14.951604 0 1
362 0.226844 -0.365416 -0.936735 -2.733887 -0.311849 0 1
401 1.018394 0.987099 0.658283 -1.609716 -0.013392 0 1
455 0.296773 0.890104 -0.123225 -0.543750 -0.053373 0 1
601 0.299993 -0.359155 -0.958818 -2.735623 -0.314648 0 1


[5 rows x 32 columns]
```

```
# Compare with Actual Fraudulent Transactions
fraud_data = data[data['Class'] == 1]
common_cases = fraud_data[fraud_data['Anomaly'] == 1]

print(f"Fraud Cases Detected by Isolation Forest: {len(common_cases)}")

# Visualize Anomalies
sns.scatterplot(data=data, x='Amount', y='Time', hue='Anomaly', palette='Set1')
plt.title("Fraud Detection (Isolation Forest)")
plt.show()
```


```
→ Fraud Cases Detected by Isolation Forest: 289
```




## ✓ Task 4: Model Evaluation

```
from sklearn.metrics import confusion_matrix, roc_curve, auc
```

```
# Confusion Matrix for XGBoost
conf_matrix = confusion_matrix(y_test, y_pred_xgb)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Greens')
plt.title("XGBoost - Confusion Matrix")
plt.show()
```



```
# ROC Curve for XGBoost
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, label=f'XGBoost (AUC = {roc_auc:.2f})')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```



## ✓ Task 5: Model Explainability

```
import shap
```

```
# SHAP for XGBoost
```


```
explainer = shap.Explainer(xgb_model, X_test)
```

```
shap_values = explainer(X_test)
```

```
# Plotting SHAP Summary
```

```
shap.summary_plot(shap_values, X_test)
```

100%|=====| 170577/170589 [19:40:00:00]



```
# Feature Importance for XGBoost
```